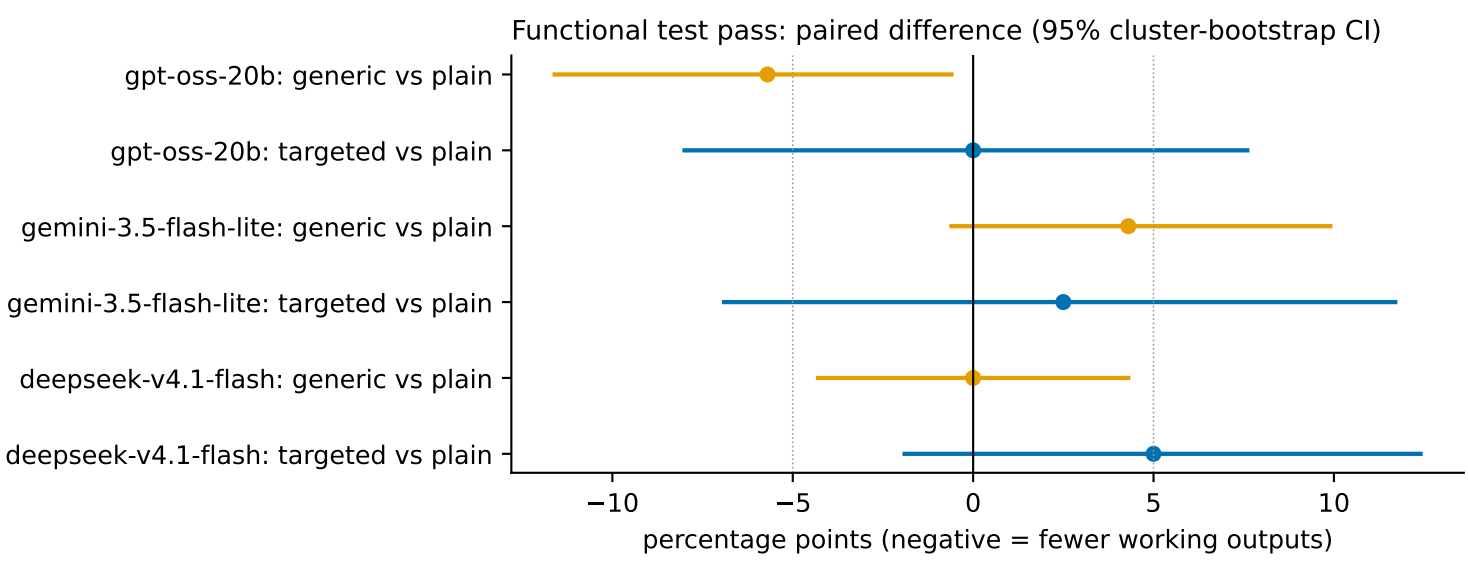
Functional test pass: paired difference (95% cluster-bootstrap CI)
gpt-oss-20b: generic vs plain
gpt-oss-20b: targeted vs plain
gemini-3.5-flash-lite: generic vs plain
gemini-3.5-flash-lite: targeted vs plain
deepseek-v4.1-flash: generic vs plain
deepseek-v4.1-flash: targeted vs plain
−10
−5
0
5
10
percentage points (negative = fewer working outputs)